FlowForge — Modern Data Orchestration Platform

- Self-service data at run in your cloud
- Build production-grade pipelines without heav ng
- AI-assisted schema detection and configuration
- Vendor-ne with your existing stack
- Simple to roll out and operat controls

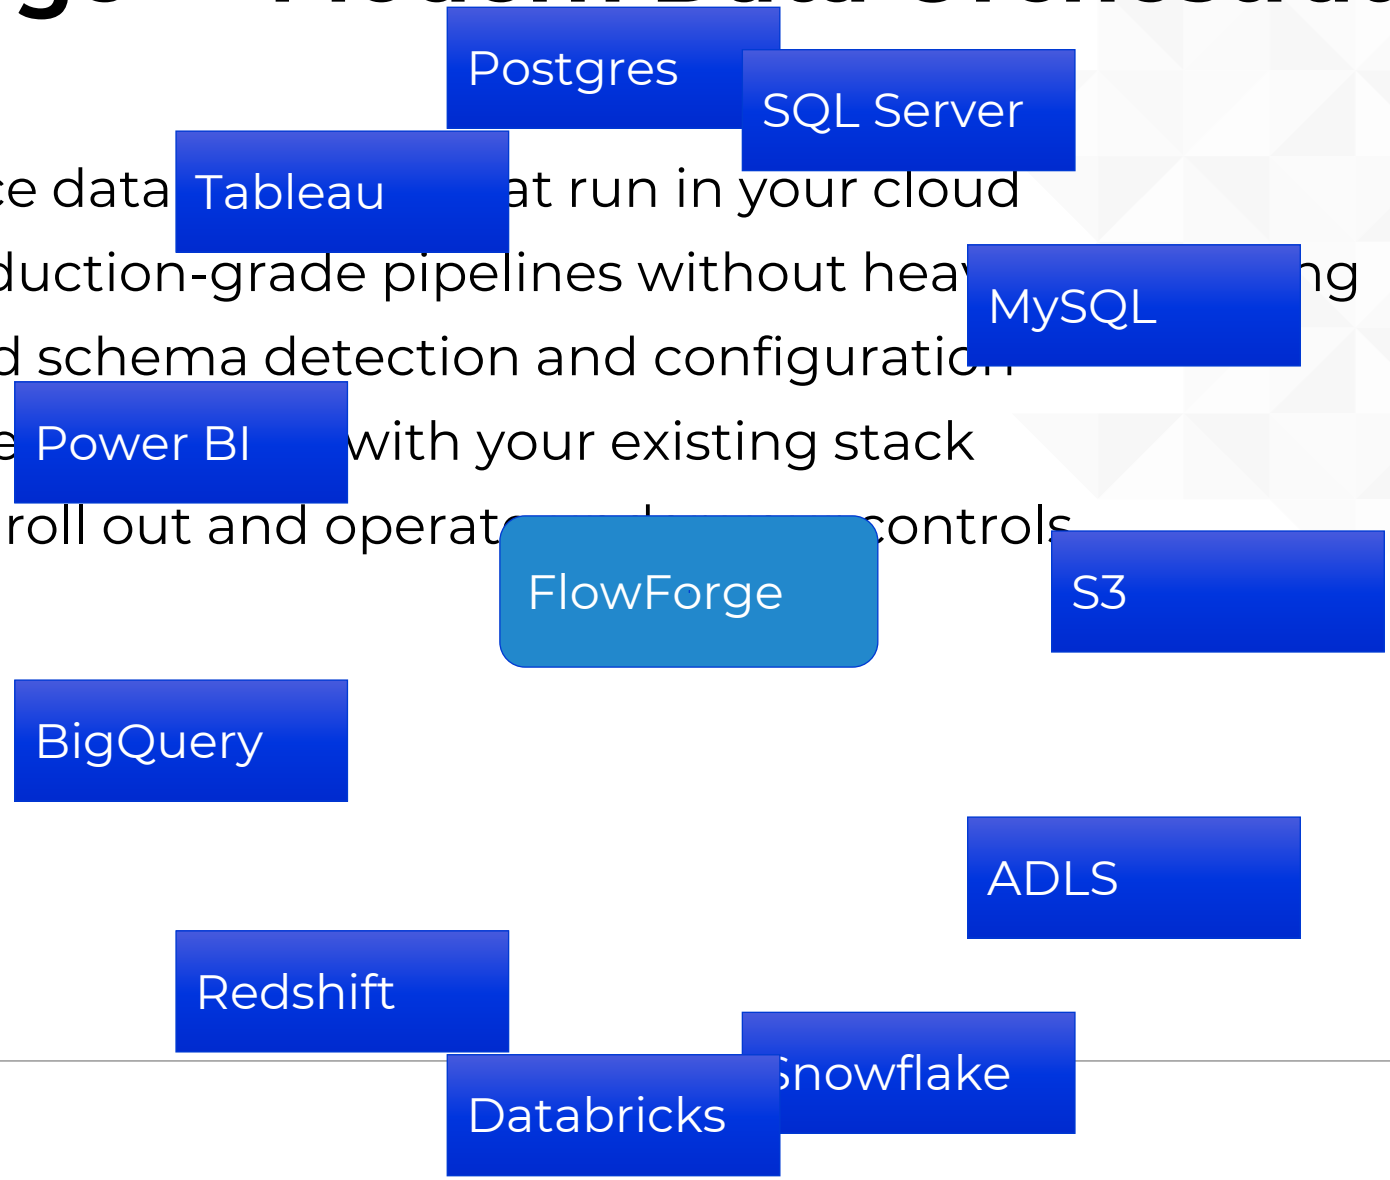

Why Data Plumbing Slows Business

- Tool sprawl and engineering bottlenecks stall delivery
- DIY code is hard to govern, maintain, and scale
- Analysts wait; leaders question data consistency
- Compliance and data residency add friction

The Case for a Simpler, Open Approach

- Cloud-hosted yet cloud-agnostic (AWS/Azure/GCP)
- Standard patterns (Medallion) with built-in governance
- Lighter operations; lower cognitive and cost overhead
- Fit with current tools; no rip-and-replace

What Is FlowForge (In Your Cloud)

- Build, run, and govern pipelines without code
- AI-powered configuration and schema detection
- Environment/team isolation and RBAC
- Built-in metadata catalog and monitoring

AI

Medallion

Governance

How It Works — Bronze → Silver → Gold

- Bronze: land and standardize files with audit columns
- Silver: cleanse, deduplicate, apply mappings/PKs
- Gold: publish compressed, analytics-ready datasets; catalog
- Triggers: manual, scheduled, and dependency-based

Six Pillars of FlowForge

- AI-assisted setup: schemas, column names, PK suggestions
- File processing at scale: pattern matching and multi-file automation
- Orchestration: multi-env, team isolation, dependency triggers
- Monitoring: live status, logs, metrics per workflow/job
- Catalog: automatic metadata, schema, lineage
- Vendor-neutral I/O: S3/ADLS, Snowflake, Databricks, Redshift, BigQuery

Why We Win

- Self-service for power users; engineers focus on complex work
- Runs in your cloud; data stays in-account for compliance
- Neutral by design; keep choice of lake/warehouse
- Predictable commercial model; avoid per-row surprises

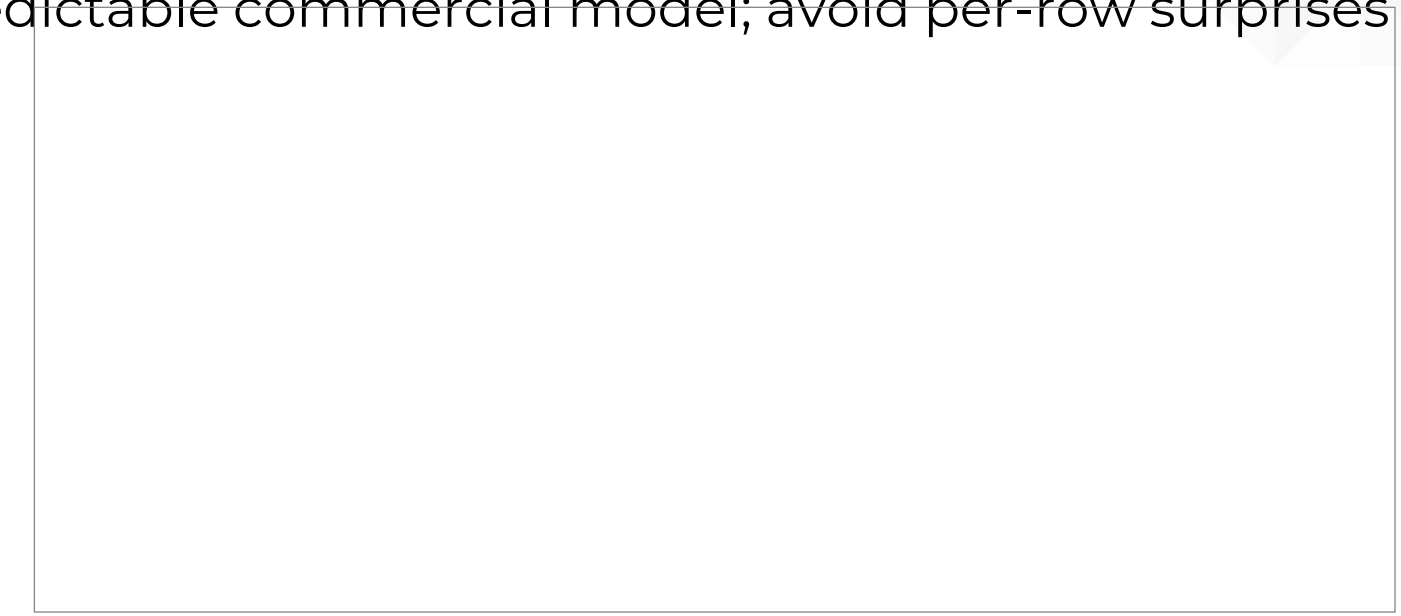

Security & Deployment (Your Cloud)

- Data residency: artifacts live in your S3/ADLS
- Identity: IAM/Key Vault integration; role-based access and approvals
- Encryption & networking: KMS/CMK, private networking, logging, audits
- Works with your monitoring and ticketing tools

What Good Looks Like (Outcomes)

- 50–80% less effort for common pipelines
- 30–60% of routine requests shift to self-service
- Consolidate tools; meaningfully reduce annual spend
- Consistent patterns increase data trust and reuse

Works With Your Stack (Ecosystem)

- Storage: Amazon S3, Azure Blob/ADLS
- Warehouses/Lakes: Snowflake, Databricks, Redshift, BigQuery
- Databases: PostgreSQL, SQL Server, MySQL (extensible)
- Analytics: Power BI, Tableau, Looker via standards-based outputs

Predictable and Simple (Packaging)

- Predictable tiers; no per-row billing complexity
- Runs in your cloud — you control infra costs
- Licensing aligned to environments/teams, not seats
- Optional low-lift pilot to validate fit

What’s Next (Highlights; No Dates)

- Additional database and API connectors
- Data quality rules and monitoring
- Document/semi-structured data processing
- Guardrails, approvals, and policy automation
- Optional streaming integrations as needed

See It on Your Data (Low-Lift Pilot)

- Bring one or two representative pipelines
- Run in your cloud with your security controls
- Keep outputs in your storage for immediate value
- Clear success criteria; minimal overhead